

FL CK: Localizing Faults in Alloy Models

Guolong Zheng, ThanhVu Nguyen, Simón Gutiérrez Brida[†], Germán Regis[‡],

Marcelo F. Frias[‡], Nazareno Luján[†], Hamid Bagheri

University of Nebraska-Lincoln

gzheng, tnguyen}@cse.unl.edu, bagheri@unl.edu

[†]University of Rio Cuarto and CONICET

sgutierrez, gregis, najuan@dc.exa.unrc.edu.ar

[‡]Buenos Aires Institute of Technology and CONICET

mfrias@itba.edu.ar

Abstract—Fault localization can help developers identify buggy statements or expressions in programs. Existing fault localization techniques are often designed for imperative programs (e.g., C and Java) and rely on tests to compare correct and incorrect execution traces to identify suspicious statements. In this demo paper, we present FL CK, a tool to automatically locate faults for models written in Alloy, a declarative language where the models are not executed but instead converted into a logical formula and solved using a SAT solver. FL CK takes as input an Alloy model that violates some assertions and returns a ranked list of suspicious expressions contributing to the violation. The key idea is to analyze the differences between *counterexamples*, i.e., instances of the model that do not satisfy the assertion and instances that do satisfy the assertion to find suspicious expressions in the input model. An experiment with 157 Alloy models with various bugs shows the efficiency and accuracy of FL CK in localizing the causes of these bugs. FL CK and its evaluation benchmark and results can be downloaded from <https://github.com/guolong-zheng/flack>. The video demonstration is available at <https://youtu.be/FKa2ohqIUms>.

Index Terms—Alloy, fault localization, specifications, models

I. INTRODUCTION

The Alloy specification language [1] has been used for various software modeling and analysis tasks such as program verification [2], test case generation [3], network security [4], and security analysis of IoT and Android platforms [5], [6]. Similar to developing programs in an imperative language like C or Java, developers can make subtle mistakes when using Alloy in modeling system specifications, especially those that capture complex systems with non-trivial behaviors. However, traditional fault localization techniques for imperative programs do not directly apply to a specification language such as Alloy, in which there are no control flow graphs or program execution traces often used to aid debugging.

To aid developers in debugging Alloy models, in [7] we introduce the FL CK technique for automatically localizing Alloy buggy expressions causing assertion violations. Given an Alloy model and properties specified by assertions, FL CK first queries the Alloy analyzer for a counterexample, an instance of the model that does not satisfy the property. Next, FL CK uses a partial max sat (PMXSAT) solver to find an instance that does satisfy the property and is as close as possible to the counterexample. FL CK then determines the relations and atoms that are different between the counterexample and

the satisfying instance. These differences explain how the counterexample violates the assertion. Finally, FL CK analyzes these differences to identify possible buggy expressions. Experimental evaluation of FL CK on 157 Alloy models with a wide variety of bugs shows that FL CK efficiently and consistently ranks buggy expressions in the top 2% of the suspicious list results.

FL CK is different than AlloyFL [8]—the only other Alloy fault localization technique currently available—in that AlloyFL relies on unconventional unit tests while FL CK uses assertions that are natural in the development practices in Alloy. Also, instead of statistically analyzing the effects of tests as in AlloyFL, FL CK relies on counterexample and satisfying instances generated by constraint solving, which are the main underlying technology in Alloy.

In this paper, we focus on the demonstration, implementation, and usage of FL CK. FL CK is highly automatic: the user only needs to provide an assertion to specify the expected behaviors, and FL CK automatically analyzes potential assertion violations and returns a ranked list of expressions based on their suspicious level to the violations. The source code and dataset of FL CK are publicly available at [9]. The full details of FL CK are available in the research paper [7].

II. THE ALLOY ANALYZER

Alloy [1] is a declarative language based on first-order logic, with an analysis engine that relies on a SAT solver. To check that an Alloy model conforms to given assertions, the Alloy analyzer automatically converts the model and assertions into a boolean formula and uses a SAT solver to search for potential counterexamples violating the assertions.

We now introduce key Alloy terminologies and concepts using the address book model in Figure 1. This model first declares three types (`sig`) `Address`, `Name` and `Book`. The type `Book` has two fields `entry` and `listed`, where `entry` is a set of `Name` and `listed` maps `entry` to a set of `Address` and `Name`.

On lines 8–10, the model defines the function (`fun`) `lookup`, which finds all `Address` and `Name` associated with a `Name` in a `Book`. Next, on lines 11–15, the model has a `fact` constraint that specifies that each `Book`’s `entry` maps to at most one `Name` or `Address`. The user makes a mistake

```

1  bstr ct sig Listing { }
2  sig ddrress extends Listing { }
3  sig Name extends Listing { }
4  sig Book {
5    entry: set Name,
6    listed: entry ->set Listing
7  }
8  fun lookup [b:Book, n:Name] : set Listing {
9    n.^(b.listed)
10 }
11 f ct {
12   ll b:Book | ll n:b.entry |
13 // Fix: replace "lone" with "some".
14   lone b.listed[n]
15 }
16 f ct {
17   ll b:Book | ll n,l:Name |
18   l in lookup[b,n] implies
19   l in b.entry
20 }
21 f ct {
22   ll b:Book | ll n:b.entry |
23   not n in lookup[b,n]
24 }
25 f ct {
26   one Book nd one ddrress nd #Name = 2
27 }
28
29 ssert assert_1 {
30   ll b:Book | ll n:b.entry |
31   some (lookup[b,n]& ddrress)
32 }
33 check assert_1
34
35 pred pred_1 {
36   ll b:Book | ll n:b.entry |
37   some (lookup[b,n]& ddrress)
38 }
39 run pred_1

```

Fig. 1: Buggy ddrress Book Model

on line 14 that uses `lone` (at most one) instead of `some` (at least one), which violates the assertion on lines 29– 32. The fact on lines 16– 20 specifies that all names reachable from any name entry in the book are themselves entries; the fact on lines 21– 24 specifies that name entries are acyclic; and the fact on lines 25– 27 specifies that there should be exactly one Book, one ddrress and two Names.

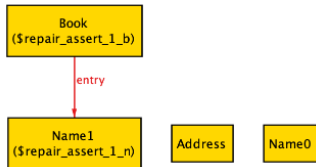


Fig. 2: counterexample generated by the running check command 33.

Iloy uses `assert` to specify assertions and the `check` command to search for counterexamples violating the asserted

properties. For example, the assertion on lines 29– 32 specifies that all name entries map to at least one address. This assertion does not always hold, and thus the check on line 33 can generate counterexamples showing its violation, e.g., the counterexample in Figure 2 shows an instance of the model in which `Name1` does not map to any ddrress. This violation is caused by the overconstraint on line 14 that uses `lone` (at most one) instead of `some` (at least one). In the next section, we demonstrate how to use FL CK to locate the buggy expression causing this violation.

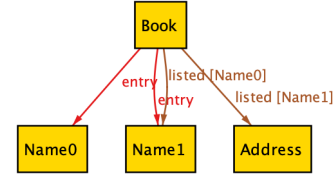


Fig. 3: satisfying instance.

Iloy uses `pred` to specify predicates and the `run` command to find instances satisfying the predicate properties. For example, the predication on lines 35– 38 specifies the same property that all name entries map to at least one address, and the `run` on line 39 can generate instances that satisfy this property (e.g., the instance in Figure 3 shows an instance in which `Name1` maps to ddrress).

III. FL CK

The FL CK tool builds on top of Iloy 4.2 and is implemented in about 8k LOC in Java. FL CK takes as input an Iloy model with assertions and returns a ranked list of suspicious expressions contributing to the violations. The key insight is that the differences between counterexamples and closely related satisfying instances can help identify expressions causing the violation in the input model. To achieve this, FL CK uses a specialized S T solver to find satisfying instances that are *as close as possible* to the counterexamples (satisfying instances generated by Iloy analyzer are random and can be vastly different than the counterexample, e.g., the satisfying instance in Figure 3 and the counterexample in Figure 2). Next, FL CK analyzes the differences between the counterexamples and the satisfying instances to find expressions in the model that likely cause the errors. Finally, FL CK computes and returns a ranked list of suspicious expressions.

. Implementation and Demonstration

Figure 4 gives an overview of the FL CK implementation, which consists of four main phases described below.

Instance generation. This phase is used to obtain pairs of counterexamples and closely similar satisfying instances. FL CK analyzes and compares the differences among these instances to identify likely buggy expressions in the model.

To generate satisfying instances similar to counterexamples, we replace the Iloy’s backend S T solver Kodkod [10] with Pardinus [11], a PM XS T (Partial M Ximum S T-ifiability) solver building on top of Kodkod. We set the



Fig. 4: Workflow of FL CK.

lloy specification with satisfying predicates as the hard constraints and the counterexamples as the soft constraints. The PM XS T solver finds a solution that satisfies all hard constraints and maximum number of soft constraints, which results in an instance that satisfies the properties specified by the assertions and most similar to the counterexample.

For the address book model, given the counterexample in Figure 2, FL CK generates the satisfying instance in Figure 5. Notice that this instance is similar to the counterexample in Figure 2, but has an extra edge from Book to ddress labeled with listed[Name1].

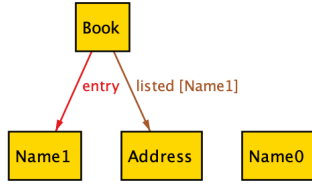


Fig. 5: satisfying instance that is close to the counterexample in Fig 2.

Difference nalysis. Given pairs of counterexamples and satisfying instances, this phase finds the minimal differences between each pair by comparing the atoms, tuples and relations of each pair. The minimal differences represent the minimal changes needed to convert a counterexample to a satisfying instance, which provide essential information related to the assertion violation.

For the instance pair in Figure 2 and Figure 5, FL CK finds the difference listed:Book->Name1-> ddress, i.e., the counterexample does not contain the tuple Book->Name1-> ddress in relation listed. In other words, the assertion violation is most related to the relation listed and three atoms: Book, Name1 and ddress. FL CK uses this information to identify the errors in the following steps.

Slicing. This phase selects expressions most related to the error by collecting expressions that contain the related relations and filtering out other expressions that only contain unrelated relations.

In the address book example, all expressions in fun and facts on lines 11– 24 are collected, as they all contain the related relation listed. The expressions in fact on lines 25– 27 are sliced out as it does not contain the relation listed, thus are not closely related to the error.

Ranking. The step computes a suspicious score for all collected expressions and their subexpressions and outputs a ranking list, as shown in Figure 6. For each expression, FL CK computes the score by iteratively calculating the scores of all

the subexpressions. The score is computed by first instantiating each expression using the values from the differences, then evaluating the instantiated expressions in both counterexamples and satisfying instances, and finally computing the score based on the differences between the evaluated results. Finally, FL CK outputs a ranking list of all expressions and their subexpressions based on the calculated scores.

It is worth noting that our implementation of FL CK does not use the standard lloy’s ST, which is designed for generating instances and is hard to modify. Instead, we extend the original lloy ST with visitors to traverse and modify ST nodes, so that we can analyze and instrument expressions with concrete values (as used in the Slicing and Ranking phases above).

IV. US GE

FL CK can be used through the command line on operating systems that supports Java. The lloy nalyzer itself is also a Java application; thus, FL CK and the lloy nalyzer can be used in the same development environment.

We demonstrate FL CK’s usage using the address book example in Figure 1. In this example, the user writes the address book model, with a mistake on line 14 that uses lone instead of some. The user also inserts an assertion and predicate on lines 29– 39 to specify the expected behaviors.

In the root directory of FL CK, the user asks FL CK to locate errors using the command in Listing 1, where the -cp ./libs/*:flack.jar option specifies the path to all dependent libraries and the jar file of flack, the option -f fsm.als tells FL CK the path to the lloy model, and the option -m 5 specifies the maximum number of pairs of counterexamples and satisfying instances to generate.

```

1 j v -Djava.library.path=solvers -cp
  ./libs/ :flack.jar loc -f addr.als -m 5
  
```

Listing 1: Command to run FL CK

Results: For this example, FL CK runs in 0.84 seconds on a 2.2 GHz Intel Core i7 CPU with 16 GB memory, and successfully points to the lone error by ranking it first among the suspicious expressions. FL CK outputs a ranking list, as shown in Figure 6 ¹. The R NK LIST is the list of suspicious expressions ranked based on their suspicious score. Each line in the ranking list consists of the ranking position and an expression followed by a suspicious score calculated by FL CK. For example, the first line in the ranking list shows that FL CK correctly ranked lone ((n.(b.listed))) ² first with a suspicious score of 1.31.

¹Due to randomness of lloy analyzer and the back-end solver, the result may be different.

²This is the syntax desugar of lone b.listed[n]

```

/flack/benchmark/addr.als:
example generation time:0.525
R NK LIST:
0: lone ((n . (b . listed))) 1.31
1: n in lookup[b,n] 1.30
2: (n . ^((b . listed))) 1.30
3: l in lookup[b,n] 1.30
4: !(n in lookup[b,n]) 1.30
5: l in (b . entry) 1.17
6: l in lookup[b,n] => l in (b.entry) 1.00
-----
analyze time(sec): 0.84
# rel: 1
# val: 3
# Slice Out: 10
# Total ST: 74
LOC: 21
evals: 368
=====

```

Fig. 6: FL CK’s results obtained for the model in Figure 1.

Additional Output: FL CK also outputs additional information about the execution for further analysis. The `example generation time` in Figure 6 records the time in seconds used by the Alloy analyzer and the PM XS T solver to generate all counterexamples and closely similar instances. The `analyze time(sec)` records the total runtime of FL CK.

The `# rel` and `# val` are the average numbers of different relations and different values among all counterexamples and satisfying instances, respectively. In this address book model, there are averagely one relation and three atom values that are different between counterexamples and satisfying instances. The `# Slice Out` and `# Total ST` are the number of sliced out and total ST nodes. In this example, FL CK sliced 10 ST nodes out of 74 nodes. The `LOC` is a rough calculation of the lines of code in the model.

The `evals` records the total number of expressions instantiated by the different values. In this example, FL CK evaluated a total number of 368 instantiated expressions.

Note that FL CK is highly automated and uses only one main parameter `-m` to specify the number of pairs of counterexamples and satisfying instances generated. Increasing the number usually results in a more accurate ranking result with a longer execution time. In the usage example, FL CK successfully ranks the buggy expression `lone n . (b . listed)` first using the option `-m 5`. However, if we use `-m 1` for this example, an expression unrelated to the error `l in (b . entry)` is ranked first with a runtime of 0.53 seconds.

V. EXPERIMENT

We evaluated FL CK on a benchmark consisting of 157 Alloy models with different kinds of bugs [7]. These include the 152 Alloy models with real faults used in the AlloyFL work [8] and 5 other Alloy models used in complex and real-world applications (e.g., surgical robots [12], android

permissions models [6], and Java program modeling and verification [2]).

The experiment results show that FL CK is able to consistently rank buggy expressions in the top 1.9%(average) of the suspicious list. FL CK successfully rank the buggy expressions for 147 out of 157 models, with 91 (58%) ranked in top 1, 38 (24%) ranked top 2 to 5, 10 (6.5%) ranked top 6 to 10, 8 (5%) ranked above top 10 for 6 and 10 (6.5%) not in the ranking list. For most of the models, FL CK finishes the execution under a second, giving the user instant feedback about the errors. The experiment results show that FL CK is general (can accurately locate a wide variety of bugs) and scalable (can handle complex, real-world Alloy models). The experiment details are given in [7] and publicly available at [9].

VI. CONCLUSION

We present the implementation details and usage of FL CK, a fault localization tool for Alloy based on assertions. The key idea of FL CK is to compute counterexample and satisfying instances of the model and compare their differences to locate errors. FL CK is fully automated and the user only needs to provide assertions to specify expected behaviors. The source code of FL CK, its benchmark models, and experimental results are publicly available at [9].

REFERENCES

- [1] D. Jackson, “Alloy: lightweight object modelling notation,” *CM Trans. Softw. Eng. Methodol.*, 2002.
- [2] J. P. Galeotti, N. Rosner, C. G. López Pombo, and M. F. Frias, “Analysis of invariants for efficient bounded verification,” in *International Symposium on Software Testing and Analysis*, 2010.
- [3] P. Bad, N. Guirre, V. S. Bengolea, D. Ciolek, M. F. Frias, J. P. Galeotti, T. Maibaum, M. M. Moscato, N. Rosner, and I. Vissani, “Improving test generation under rich contracts by tight bounds and incremental ST solving,” in *International Conference on Software Testing, Verification and Validation*, 2013.
- [4] F. J. Maldonado-Lopez, J. Chavarriaga, and Y. Donoso, “Detecting network policy conflicts using alloy,” in *International Conference on Abstract State Machines*, 2014.
- [5] M. Ihanahnah, C. Stevens, and H. Bagheri, “Scalable analysis of interaction threats in iot systems,” in *International Symposium on Software Testing and Analysis*, 2020.
- [6] H. Bagheri, S. Sadeghi, J. Garcia, and S. Malek, “Covert: Compositional analysis of android inter-app permission leakage,” *IEEE Transactions on Software Engineering*, 2015.
- [7] G. Zheng, T. Nguyen, S. Gutiérrez-Brida, G. Regis, M. Frias, N. Guirre, and H. Bagheri, “Flack: Counterexample-guided fault localization for alloy models,” in *International Conference on Software Engineering*, 2021.
- [8] K. Wang, J. Sullivan, D. Marinov, and S. Khurshid, “Fault Localization for Declarative Models in Alloy,” in *International Symposium on Software Reliability Engineering*, 2020.
- [9] *FL CK repository*, 2020. [Online]. Available: <https://doi.org/10.6084/m9.figshare.13439894.v4>
- [10] E. Torlak and D. Jackson, “Kodkod: relational model finder,” in *Tools and Algorithms for the Construction and Analysis of Systems*, 2007.
- [11] J. Cunha, N. Macedo, and T. Guimarães, “Target oriented relational model finding,” in *Fundamental Approaches to Software Engineering*, 2014.
- [12] N. Mansoor, J. J. Saddler, B. Silva, H. Bagheri, M. B. Cohen, and S. Farritor, “Modeling and testing a family of surgical robots: an experience report,” in *Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, 2018.